

Ensemble Determinization in Monte Carlo Tree Search for the Imperfect Information Card Game *Magic: The Gathering*

Peter I. Cowling, *Member, IEEE*, Colin D. Ward, *Student Member, IEEE*, and Edward J. Powley, *Member, IEEE*

Abstract—In this paper, we examine the use of Monte Carlo tree search (MCTS) for a variant of one of the most popular and profitable games in the world: the card game *Magic: The Gathering* (*M:TG*). The game tree for *M:TG* has a range of distinctive features, which we discuss here; it has incomplete information through the opponent's hidden cards and randomness through card drawing from a shuffled deck. We investigate a wide range of approaches that use determinization, where all hidden and random information is assumed known to all players, alongside MCTS. We consider a number of variations to the rollout strategy using a range of levels of sophistication and expert knowledge, and decaying reward to encourage play urgency. We examine the effect of utilizing various pruning strategies in order to increase the information gained from each determinization, alongside methods that increase the relevance of random choices. Additionally, we deconstruct the move generation procedure into a binary yes/no decision tree and apply MCTS to this finer grained decision process. We compare our modifications to a basic MCTS approach for *M:TG* using fixed decks, and show that significant improvements in playing strength can be obtained.

Index Terms—Card games, determinization, imperfect information, *Magic: The Gathering* (*M:TG*), Monte Carlo tree search (MCTS), parallelization.

I. INTRODUCTION

MONTE CARLO TREE SEARCH (MCTS) has, in recent years, provided a breakthrough in creating AI agents for games [1]. It has shown remarkable success in *Go* [2]–[4] and is being applied successfully to a wide variety of game environments [5], including *Hex* [6], *Havannah* [7], and general game playing [8], [9]. One of the major strengths of MCTS is that there is no requirement for a strong evaluation function, and it has therefore been especially useful for games where an evaluation function is difficult to formulate, such as *Go* [10] and *Hex* [6]. In 2000, Schaeffer [11] said “it will take many decades of research and development before world-championship-caliber *Go* programs exist” and yet we have recently seen MCTS-based *Go* players begin to challenge the best human players in the

world [2]. The lack of a requirement for any specific domain knowledge has also helped MCTS to become very successful in the area of general game playing where there is little advance knowledge of the structure of the problem and, therefore, a very restricted scope in which to develop an evaluation function [9].

Removing the need to have a sophisticated evaluation function suggests the possibility of developing search-based AI game agents for much more complex games than was previously possible, and suggests an avenue for a new AI approach in video games. The video games industry is a huge and growing market: in 2009, the video game industry had sales of over \$10 billion in the United States alone [12], and while the graphics and visual appeal of games has progressed enormously in recent years, to the extent of mapping recognizable faces and emotional content [13], the AI being used is still largely the nonadaptive scripting approach that has always been used [14].

While MCTS has made great strides in producing strong players for perfect information games, the situation for imperfect information games is less advanced and often the use of MCTS is restricted to perfect information versions or parts of the game. For example, in *Settlers of Catan* [15], the authors reduced the game to a perfect information variant and then applied MCTS to this perfect information system, convincingly beating hand-coded AI from an open source version of the game with only 10 000 simulations per turn and a small amount of domain knowledge. Perfect information variants of *Spades* and *Hearts* card games have also been used to study convergence properties of a variant of MCTS in a multiplayer environment [16].

Card games typically have a wealth of hidden information and provide an interesting challenge for AI. Chance actions covering all possible cards that may be drawn from a deck of cards yield a game tree with a chance node at the root which explodes combinatorially, quickly generating branching factors which may dwarf that of *Go*. We must also deal with the effect of hidden information, e.g., the particular cards in an opponent's unseen hand. However, card and board games offer an important class of difficult decision problems for AI research, having features in common with perfect information games and video games, and a complexity somewhere between the two.

MCTS has been applied to several card games with some success. MCTS-based players for *Poker* have started to challenge the best humans in heads-up play [17]. Advances have also been made in multiplayer *Poker* and *Skat* [18], which show promise toward challenging the best human players.

Manuscript received August 03, 2011; revised December 18, 2011; accepted June 08, 2012. Date of publication June 14, 2012; date of current version December 11, 2012. This work was supported by the U.K. Engineering and Physical Sciences Research Council (EPSRC).

P. I. Cowling and E. J. Powley are with the Department of Computer Science, University of York, York YO10 5GH, U.K. (e-mail: peter.cowling@york.ac.uk; e.powley@bradford.ac.uk).

C. D. Ward is with the Artificial Intelligence Research Centre, School of Computing, Informatics and Media, University of Bradford, Bradford, West Yorkshire BD7 1DP, U.K. (e-mail: c.d.ward@student.bradford.ac.uk).

Digital Object Identifier 10.1109/TCIAIG.2012.2204883

Determinization, where all hidden and random information is assumed known by all players, allows recent advances in MCTS to be applied to games with incomplete information and randomness. The determinization approach is not perfect: as discussed by Frank and Basin [19], it does not handle situations where different (indistinguishable) game states suggest different optimal moves, nor situations where the opponent's influence makes certain game states more likely to occur than others. In spite of these problems, determinization has been applied successfully to several games. An MCTS-based AI agent which uses determinization has been developed that plays *Klondike Solitaire* [20], arguably one of the most popular computer games in the world. For the variant of the game considered, the performance of MCTS in this case exceeds human performance by a substantial margin. A determinized Monte Carlo approach to *Bridge* [21], which uses Monte Carlo simulations with a tree of depth one has also yielded strong play. The combination of MCTS and determinization is discussed in more detail in Section V.

In this paper, we investigate MCTS approaches for the card game *Magic: The Gathering (M:TG)* [22]. *M:TG* is a strategic card game for two players, which shares characteristics with many other card games: hidden information in the opponent's hand and the stochastic nature of drawing cards from a shuffled deck. Where *M:TG* differs from other card games is that it does not use a standard deck of cards but rather cards that have been created specifically for the game. Many cards change the rules of the game in subtle ways and the interaction between the rules changes on the cards gives rise to very rich game play.

M:TG is played by over 12 million people worldwide and in 2005 the manufacturer Hasbro reported that it was their biggest selling game, outstripping *Monopoly*, *Trivial Pursuit*, and *Cluedo* [23]. The game has a number of professional players: in 2011, the professional tour paid out almost \$1 million in prize money to the best players in the world. While specific sales figures are unavailable, it is estimated that more than \$100 million is spent annually on the game [24].

M:TG is not only played with physical cards. In 2009, a version of the game appeared on Xbox Live Arcade that allowed players to play against a computer opponent. The details are proprietary but the game appears to use a depth-limited decision tree with static evaluation of the game state at leaf nodes [25]. The AI in the game has generally been regarded as plausible for someone who is a beginner to the game but is at a level that would not challenge an average player [26].

M:TG possesses several characteristics that we believe make it an interesting area for research into game AI.

- 1) *M:TG* does not use a standard deck of cards but instead uses cards that are specifically designed for the game. Players are free to construct their own deck using these cards, a decision problem of enormous complexity. There are currently over 9000 different cards that have been created for *M:TG* and more are added every year. This makes it particularly difficult to predict what cards an opponent may have in his deck and the consequent interactions between cards. It also makes *M:TG* arguably an exercise in general game playing and a step toward understanding generalizable approaches to intelligent game play.

- 2) Players are not limited to playing a single card on their turn. All cards have costs and, as the game progresses, the resources, and hence options available to a player, increase. A player may play any number of cards from their hand on their turn providing they can pay all the associated costs. This means that at each turn, a player can play a subset of cards in hand, giving a high branching factor.
- 3) The interaction between the players is highly significant, and there is substantial scope for opponent modeling and inference as to the cards the opponent holds in his hand and his deck. Inference is a critical skill in games between human players.
- 4) The sequence of play is not linear, and the opponent can "interrupt" the player's turn, for example, to cancel the effect of playing a particular card. Hence, *M:TG* is less rigid than most turn-based games as each player may have decision points during the opponent's turn.

The structure of this paper is as follows. In Section II, we discuss MCTS. In Section III, we describe the game of *M:TG* and the simplified variant of the game that we have used in our trials. In Section IV, we describe the rules-based players we have devised as opponents for our MCTS players. Section V surveys work on the use of parallel determinization approaches to handle uncertainty and incomplete information. Our enhancements to MCTS for *M:TG* which use parallel determinization are presented in Section VI. Section VII presents experimental results and analysis. Section VIII draws conclusions and provides suggestions for future work.

II. MONTE CARLO TREE SEARCH

MCTS extends ideas of bandit-based planning [27] to search trees. In the k -armed bandit problem, Auer *et al.* [27] showed that it was possible to achieve best possible logarithmic regret by selecting the arm that maximized the upper confidence bound (UCB)

$$\bar{x}_j + \sqrt{\frac{2 \ln n}{n_j}}$$

where \bar{x}_j is the average reward from arm j , n_j is the number of times arm j has been played so far, and n is the total number of plays so far.

Around 2006–2007, several teams of researchers were investigating the application of Monte Carlo approaches to trees: Chaslot *et al.* [28] developed the idea of objective Monte Carlo that automatically tuned the ratio between exploration and exploitation based on the results of Monte Carlo simulations at leaf nodes in a minimax tree. Coulom [29] described a method of incrementally growing a tree based on the outcome of simulations at leaf nodes and utilizing the reward from the simulated games to bias the tree growth down promising lines of play. Kocsis and Szepesvári used the UCB formula recursively as the tree was searched [30]. The resulting algorithm is known as UCT (UCB applied to trees), and Kocsis and Szepesvári showed that with only very limited conditions, it would produce optimal play given a very large number of simulations. The range of algorithms which use Monte Carlo simulations as an approach




Fig. 1. The four steps of an MCTS algorithm.

to heuristically building a search tree have become commonly known as MCTS.

MCTS algorithms are generally a four-step process that is repeated until some limit is reached, usually a limit on elapsed time or number of simulations. In this paper, we have commonly used a total number of simulations as the limit in our experiments.

The steps of the algorithm, illustrated in Fig. 1, are as follows.

- 1) *Selection*. The algorithm uses the UCB formula (or some other approach) to select a child node of the position currently being considered, repeating this process until a leaf node is reached. Selection balances the exploitation of known good nodes with the exploration of nodes whose value is currently uncertain.
- 2) *Expansion*. One or more children are added to the leaf node reached in the selection step.
- 3) *Simulation (or rollout)*. A simulation is carried out from the new leaf node, using a random move generator or other approach at each step, until a terminal game state is reached.
- 4) *Backpropagation*. The reward received at the simulation step is propagated back to all nodes in the tree that were part of the selection process to update the values (e.g., number of wins/visits) in those nodes.

The algorithm has two principal advantages over conventional tree search methods such as minimax with alpha-beta pruning.

- 1) It is “anytime” [31]. The algorithm can be stopped at any point to yield a result which makes use of all rollout information so far. There is no need to reach a particular stage during search before a result is obtainable, as there would be for minimax search, even with iterative deepening [32].
- 2) An evaluation function is not required for nonterminal game states, as simulation always reaches a terminal state. The reward for a given game state is obtained by aggregating win/lose simulation results from that state.

MCTS may utilize randomly selected moves when conducting simulations and therefore has no need of *any* specific domain knowledge, other than the moves available from each game state and the values of terminal game positions. In practice, however, the performance of the algorithm can usually be improved by including some domain specific considerations in the simulation and selection phases [3].

III. *MAGIC: THE GATHERING*

A. Game Rules

In the game of *M:TG* each player takes on the role of a wizard contesting a duel with their opponent. Each player’s hand of cards represents the spells and resources that the wizard has available and the players play cards from their hand in order to either generate resources or play spells with which to beat their opponent.

Each player has a life total and the player whose life total is reduced to zero first loses the game. The game consists of multiple varieties of cards and multiple types of resource, consequently the possible interactions between the available cards can become extremely complex. Much of the appeal of *M:TG* arises through understanding and tactically exploiting the interactions between the player’s cards, and between player’s and opponent’s cards.

The full game is very complex and difficult to model easily, so we have chosen to retain the basic structure and turn order mechanics of the game but to focus on creature combat, which is the most important form of interaction between *M:TG* cards for the majority of decks (and essentially for all decks played by beginning human players). By restricting the test environment to only land (resource) cards and creature (spell) cards we simplify encoding of the rules (which represents a significant software engineering problem in practice [33]). In our test version of the game the players have a deck of cards containing only creatures and land resource cards of a single color.

Each *creature* card has *power* and *toughness* values denoting how good the creature is at dealing and taking damage, respectively, and a resource (or mana) *cost*. In general, more powerful creatures have a higher resource cost. Below we will refer to a creature with power P , toughness T , and cost C as $P/T^{(C)}$, and omit C when it is not significant to our discussion. At each turn a player may put at most one *land* resource card into play from their hand, referred to below as L .

Over the course of the game, each player will accumulate land cards in play. On any given turn the player may expend resources equal to the total amount of land they have in play in order to meet the costs of creature cards from their hand. This allows them to play creature cards from their hand to the *in play zone* which are then available to *attack* and *defend*. These spent



Fig. 2. Representation of the play area during a game of *M:TG*.

resources refresh at the beginning of the player's next turn. In this way, as the player controls increasing amounts of land, they can afford more expensive creatures.

Creatures may be available to defend against the opponent's attack although they are not required to do so. Creatures that have attacked on a defending player's previous turn are considered *tapped* and therefore are not available for defence. Once attacking creatures are declared, the defending player allocates each untapped defending creature (a *blocker*) to, at most, one attacker. Each attacking creature may have none, one, or more blockers assigned to it. Blocking creatures die, and are consequently removed from play, to the *graveyard*, if the attacking creature allocates *damage* to a blocker greater than or equal to its toughness. The attacking creature dies if the corresponding blockers' total power provides damage greater than or equal to the attacking creature's toughness. In the case of an attacking creature having multiple blockers, then the player controlling the attacking creature decides how to split each attacker's damage among its blockers. Creatures that are not blocked cause damage to the opponent's *life total* and a player loses the game if his life total is zero or less. See Fig. 2.

B. Structure of a Game Turn

The players take turns. On any given turn, one player is "active" and the other is "nonactive" and can merely respond to the actions of the active player. The sequence of events during a typical turn is as follows.

- 1) The active player draws a card from his deck and adds it to his hand. If he is unable to do so (because his deck has no remaining cards) then he immediately loses the game.
- 2) The active player selects a subset of his creatures in play to be attackers.
- 3) The nonactive player assigns each untapped creature he has in play to block, at most one, attacker.
- 4) Combat is resolved and any creatures taking sufficient damage are removed from play. Any unblocked attackers do damage to the nonactive player's life total. If the nonactive player's life total falls to zero, then that player loses the game.
- 5) The active player may play cards from his hand. One land card may be played each turn and the accumulated land in play can be used to pay for cards to be played from his hand provided the total cost of creatures played is less than the total number of land cards in play.
- 6) The active and nonactive players then switch roles and a new turn begins.

IV. A RULE-BASED APPROACH TO *M:TG*

Two rule-based players were created of differing play strength, as well as a purely random player, in order to provide test opponents and rollout strategies for our MCTS players. The first rule-based player had the best heuristics we were able to devise in all areas of the game, and was created using substantial human expertise. The second rule-based player had a significantly reduced set of heuristics and included elements of randomness in its decisions.

From the structure of the game turn, we can see that there are three main decisions that a player needs to make. The active player must decide which (if any) creatures will attack, and then which cards to play after the attack is resolved. The nonactive player must decide how to block the attacking creatures.

Attacking and blocking decisions in *M:TG* are far from trivial. Multiple creatures may be selected to attack and each of those attacking creatures may be blocked by any subset of the defenders creatures (subject to the constraint that a defending creature can only block one attacking creature). There are also considerations that creatures selected to attack are unavailable for defensive duty on the next turn so the attacking player has to avoid leaving himself open to a lethal counter attack. We were fortunate to be able to call on the experience of a strong *M:TG* player in order to aid us in formulating some attacking and blocking heuristics.

A. "Expert"-Rule-Based Player

Here we present the detailed description of the heuristics utilized by the expert-rule-based player. There are separate heuristics for attacking, blocking, and playing cards from the hand.

The CHOOSEATTACKERS function (Algorithm 1) decides which creatures from those available to the player will be selected to attack the opponent this turn. The basic approach taken is to consider each creature that could attack this turn and determine whether there is a reason that it should "not" attack. If no such reason is found, then the creature is declared as an attacker.

Algorithm 1 Attacker choice for the expert rule-based player (Section IV-A).

```

1: function CHOOSEATTACKERS( $P_A, P_B, l_A, l_B, m_A$ )
2:   parameters
3:      $P_A = (\text{potential attackers}) = (p_n/t_n^{(\alpha_n)}, p_{n-1}/t_{n-1}^{(\alpha_{n-1})}, \dots, p_1/t_1^{(\alpha_1)})$ 
       where  $p_n \geq p_{n-1} \geq \dots \geq p_1$  and  $p_i = p_{i-1} \implies \alpha_i \geq \alpha_{i-1}$  ( $i = 2, 3, \dots, n$ )
4:      $P_B = (\text{potential blockers}) = (q_m/s_m^{(\beta_m)}, q_{m-1}/s_{m-1}^{(\beta_{m-1})}, \dots, q_1/s_1^{(\beta_1)})$ 
       where  $q_m \geq q_{m-1} \geq \dots \geq q_1$ 
5:      $l_A, l_B =$  life total for attacking and blocking player, respectively
6:      $m_A =$  maximum number of creatures attacking player has enough land to play from his hand this turn
7:      $d = |P_A| - |P_B|$ 
8:      $a_{\max} = |P_A| + m_A - \min\{i : q_i + q_{i-1} + \dots + q_1 \geq l_A\}$ 
9:
10:    decision variables
11:     $A = \{\text{chosen attackers}\} \subseteq P_A$ 
12:
13:    // Special cases
14:    if  $P_A = \emptyset$  then return  $A = \emptyset$ 
15:    else if  $P_B = \emptyset$  then return  $A = P_A$ 
16:    else if  $d > 0$  and  $p_d + p_{d-1} + \dots + p_1 \geq l_B$  then return  $A = P_A$ 
17:    else if  $a_{\max} = 0$  then return  $A = \emptyset$ 
18:    else if  $a_{\max} < 0$  then return  $A = P_A$ 
19:    end if
20:
21:    // Main case
22:     $i \leftarrow n$ ;  $A \leftarrow \emptyset$ 
23:    do
24:      if there is no  $M \subseteq P_B$  with  $s_j > p_i$  (for all  $q_j/s_j^{(\beta_j)} \in M$ ) and  $\sum_{k \in M} q_k \geq t_i$ 
25:        and there is no pair  $(M' \subseteq P_B, q_b/s_b^{(\beta_b)} \in (P_B \setminus M'))$ 
26:          with  $\beta_b < \alpha_i$  and  $s_j > p_i$  (for all  $q_j/s_j^{(\beta_j)} \in M'$ ) and  $q_b + \sum_{k \in M'} q_k \geq t_i$ 
27:          and there is no  $q_b/s_b^{(\beta_b)} \in P_B$  with  $p_i > s_b$  and  $\alpha_i < \beta_b$ 
28:        then
29:           $A \leftarrow A \cup \{p_i/t_i^{(\alpha_i)}\}$ 
30:        end if
31:         $i \leftarrow i - 1$ 
32:    while  $|A| < a_{\max}$  and  $i > 0$ 
33:    return  $A$ 
34: end function

```

Lines 14–18 of CHOOSEATTACKERS (Algorithm 1) define special cases. If there are no potential attackers, then there will be no attackers (line 14). If there are no potential blockers, or the number of blockers is too small to prevent lethal damage (lines 15 and 16, respectively), then we attack with all potential attackers. a_{\max} defines the maximum number of attackers to leave sufficient blockers on the next turn to prevent lethal damage. If a_{\max} is zero, then we will have no attackers (line 17), and if a_{\max} is less than zero, we will lose next turn anyway, so we attack with all possible creatures to maximize the chance that the opponent might make a mistake in blocking (line 18). In the main case, we then go through possible creatures by descending power (breaking ties by descending cost) and choose to attack with a creature if there is no set of blockers that can block and kill it without any blocker being killed (line 24); no blocking combination that kills the attacker and results in only a single blocker of lower mana cost than the attacker being killed (line 25); and the attacker cannot be held back to block and kill an opposing creature of higher mana cost next turn (line 26).

The CHOOSEBLOCKERS function (Algorithm 2) is constructed by considering possible ways to block each attacking creature in descending order of attractiveness to the defending player. Ideally the attacking creature should be killed with no loss to the defender, but if this is not possible, then lesser outcomes are examined until ultimately, if the defending player must block because otherwise he will lose and no better outcome can be discovered, it will “chump” block with its weakest creature. This creature will certainly die but it will prevent damage reaching the player.

Lines 14–16 of CHOOSEBLOCKERS (Algorithm 2) define special cases. If there are no attackers or no creatures available to block, then no blockers need to be declared (lines 14 and 15). b_{\min} defines the minimum number of attacking creatures that need to be blocked in order for the defending player to survive the attack. If this is higher than the number of potential blockers, then game loss is certain and there is no point looking for blocks (line 16). In the main case, we look at each attacking creature in descending order of power (break ties by descending mana

Algorithm 2 Blocker choice for the expert rule-based player (Section IV-A).

```

1: function CHOOSEBLOCKERS( $A, P_B, l_B$ )
2:   parameters
3:      $A = (\text{chosen attackers}) = (p_k/t_k^{(\alpha_k)}, p_{k-1}/t_{k-1}^{(\alpha_{k-1})}, \dots, p_1/t_1^{(\alpha_1)})$ 
       where  $p_k \geq p_{k-1} \geq \dots \geq p_1$  and  $p_i = p_{i-1} \implies \alpha_i \geq \alpha_{i-1}$  ( $i = 2, 3, \dots, k$ )
4:      $P_B = (\text{potential blockers}) = (q_m/s_m^{(\beta_m)}, q_{m-1}/s_{m-1}^{(\beta_{m-1})}, \dots, q_1/s_1^{(\beta_1)})$ 
       where  $q_m \geq q_{m-1} \geq \dots \geq q_1$ 
5:      $l_B = \text{life total for blocking player}$ 
6:      $b_{\min} = \text{minimum number of blockers} = \begin{cases} \min \{i : p_i + p_{i-1} + \dots + p_1 \geq l_B\} & \text{if } p_k + p_{k-1} + \dots + p_1 \geq l_B \\ 0 & \text{otherwise} \end{cases}$ 
7:
8:   decision variables
9:      $B_{(i)} = \{\text{blockers chosen for attacker } p_i/t_i^{(\alpha_i)}\} \subseteq P_B$ 
10:     $\mathcal{B} = (\text{all blocks}) = (B_{(1)}, B_{(2)}, \dots, B_{(k)})$ 
11:     $B = \{\text{all blocking creatures}\} = \bigcup_i B_{(i)}$  (note  $B_{(i)} \cap B_{(j)} = \emptyset$  for  $i \neq j$ )
12:
13:  // Special cases
14:  if  $A = \emptyset$  then return  $\mathcal{B} = ()$ 
15:  else if  $P_B = \emptyset$  then return  $\mathcal{B} = (\emptyset, \emptyset, \dots, \emptyset)$ 
16:  else if  $b_{\min} > |P_B|$  then return  $\mathcal{B} = (\emptyset, \emptyset, \dots, \emptyset)$ 
17:  end if
18:
19:  // Main case
20:   $i \leftarrow k$ 
21:  do
22:     $P = P_B \setminus B$ 
23:     $\mathcal{Q} = \left\{ Q \subseteq P : s_j > p_i \text{ for all } q/s^{(\beta)} \in Q \text{ and } \sum_{q/s^{(\beta)} \in Q} q \geq t_i \right\}$ 
24:    if  $\mathcal{Q} \neq \emptyset$  then choose  $B_{(i)} \in \arg \min_{Q \in \mathcal{Q}} \sum_{q/s^{(\beta)} \in Q} \beta$ ; goto line 34
25:     $\mathcal{Q}' = \{q/s^{(\beta)} \in P : q \geq t_i \text{ and } \beta < \alpha_i\}$ 
26:    if  $\mathcal{Q}' \neq \emptyset$  then choose  $B_{(i)} \in \arg \min_{q/s^{(\beta)} \in \mathcal{Q}'} \beta$ ; goto line 34
27:     $\mathcal{Q}'' = \{(q_x/s_x^{(\beta_x)}, q_y/s_y^{(\beta_y)}) \in P^2$ 
       :  $x \neq y, \beta_x \leq \beta_y, q_x + q_y \geq t_i, s_x + s_y > p_i$  and  $\beta_j \leq \alpha_i$  if  $s_j \leq p_i$  for  $j \in \{x, y\}\}$ 
28:    if  $\mathcal{Q}'' \neq \emptyset$  then choose  $B_{(i)} \in \arg \min_{(q/s^{(\beta)}, q'/s'^{(\beta')}) \in \mathcal{Q}''} \beta$ ; goto line 34
29:     $\mathcal{Q}''' = \{q/s^{(\beta)} \in P : s > p_i\}$ 
30:    if  $\mathcal{Q}''' \neq \emptyset$  then choose  $B_{(i)} \in \arg \min_{q/s^{(\beta)} \in \mathcal{Q}'''} \beta$ ; goto line 34
31:     $\mathcal{Q}'''' = \left\{ Q \subseteq P : \sum_{q/s^{(\beta)} \in Q} q \geq t_i \right\}$ 
32:    if  $i > k - b_{\min}$  and  $\mathcal{Q}'''' \neq \emptyset$  then choose  $B_{(i)} \in \arg \min_{Q \in \mathcal{Q}''''} \sum_{q/s^{(\beta)} \in Q} \beta$ 
33:    else if  $i > k - b_{\min}$  then choose  $B_{(i)} \in \arg \min_{q/s^{(\beta)} \in P} \beta$ 
34:     $P_B \leftarrow P_B \setminus B_{(i)}$ 
35:     $i \leftarrow i - 1$ 
36:  while  $P_B \neq \emptyset$  and  $i > 0$ 
37:  return  $\mathcal{B} = (B_{(1)}, B_{(2)}, \dots, B_{(k)})$ 
38: end function

```

cost) and evaluate the best blocking option. These options are evaluated in a descending order of favorability for the defending player so that, once an option is found whose conditions are met, we greedily assign that set of blockers and move on to the next attacking creature. First, we see if there is any set of blockers that would kill the attacker without any of the blockers dying. If

such a set exists, we select the one that has the minimum total mana cost (line 24). Then, we see if there is a single creature that would kill the attacker and has a lower mana cost than the attacker (line 26); our blocking creature would die but we would lose a less valuable creature than the attacking player. We then look for a pair of blockers that together can kill the attacker

Algorithm 3 Card choice for the expert rule-based player (Section IV-A).

```

1: function CHOOSEMAIN( $L_A, C_A, m$ )
2:   parameters
3:      $L_A = \{\text{land cards in active player's hand}\}$ 
        $= \{L, L, \dots, L\}$ 
4:      $C_A = \{\text{creature cards in active player's hand}\}$ 
        $= (p_n/t_n^{(\alpha_n)}, p_{n-1}/t_{n-1}^{(\alpha_{n-1})}, \dots, p_1/t_1^{(\alpha_1)})$ 
       where  $\alpha_n \geq \alpha_{n-1} \geq \dots \geq \alpha_1$ 
5:      $m = \text{total mana available to active player}$ 
6:
7:   decision variables
8:      $P_A = \{\text{cards to play this turn}\} \subseteq L_A \cup C_A$ 
9:
10:  // Play land
11:  if  $L_A \neq \emptyset$  then
12:     $P_A \leftarrow P_A \cup \{L\}$ 
13:     $m \leftarrow m + 1$ 
14:  end if
15:
16:  // Play creatures
17:   $i \leftarrow n$ 
18:  do
19:    if  $\alpha_i \leq m$  then
20:       $P_A \leftarrow P_A \cup \{p_i/t_i^{(\alpha_i)}\}$ 
21:       $m \leftarrow m - \alpha_i$ 
22:    end if
23:     $i \leftarrow i - 1$ 
24:  while  $m > 0$  and  $i > 0$ 
25:  return  $P_A$ 
26: end function

```

while only losing one of their number with a smaller mana cost than the attacker (for example, a $4/4^{(5)}$ attacker blocked by a $2/2^{(2)}$ and a $2/3^{(3)}$), and the pair which leads to the lowest mana cost blocker being killed is chosen (line 28).

So far we have only considered blocks that are advantageous to the defending player. We then look at the neutral case where we block with a creature that will not die to the attacker but will not kill the attacker (line 30). Finally, we check whether we need to look at disadvantageous blocks. If $i > k - b_{\min}$, then we must block this attacker or the player will die. First, we find the lowest mana cost group that kills the attacker (line 32), or if no such group exists, we assign the lowest cost blocker still available to “chump” block (line 33) so avoiding lethal damage to the player.

The rules for selecting cards to play are much simpler than the attacking and blocking rules. In CHOOSEMAIN (Algorithm 3), we use a greedy approach that plays land if possible (line 11) and plays out the most expensive affordable creature in the player’s hand greedily (line 19) until the player cannot afford any more creatures.

B. “Reduced”-Rule-Based Player

The reduced-rule-based player utilizes much simpler heuristics for its decisions and includes randomness in the decision-making process. This approach is significantly weaker than the player given above but gives the possibility of choosing any

attacking/blocking move, and any nondominated move in the main phase. Our intuition suggests that this may be effective in constructing the MCTS tree and in conducting rollouts.

- **CHOOSEATTACKERS:** For each creature that is able to attack, the player decides with probability p whether to attack with that creature. For our tests $p = 0.5$, so that the player chooses uniformly at random from possible subsets of attackers.
- **CHOOSEBLOCKERS:** For each available creature that can block, the player decides uniformly at random among all the available attacking creatures plus the decision not to block anything and assigns the blocking creature accordingly.
- **CHOOSEMAIN:** This player uses the same approach to CHOOSEMAIN as the “expert”-rule-based player, but with the modification that it uses an ordering of creatures in hand, chosen uniformly at random from all orderings. Hence, any nondominated play can be generated. Here nondominated means that after the main phase cards are played, there remain no more playable cards in the active player’s hand.

We ran a direct comparison between our two rule-based players in order to gauge their relative strength. We ran an experiment of 1000 randomly generated test games ten times (playing 10 000 games in total) in order to generate confidence interval information. The expert-rule-based player proved to be much stronger, winning 63.7% of games with a 95% confidence interval of $\pm 0.94\%$.

C. Performance Against Human Opponents

We also tested the ability of the expert-rule-based player against a number of human opponents. A total of 114 games were played against seven human players. Six of the human players rated themselves as strong—winning at local events and competitive within the regional/national scene; one player considered himself as a little less strong, rating himself competitive at local events. All the human players played between 10 and 25 games against the expert-rule-based player.

Overall the expert-rule-based player won 48 of the 114 games played for a win rate of 42.1%. The expert-rule-based player performed slightly better when playing first in a game and won 27 out of 58 games for a win rate of 46.6%. The expert-rule-based player performed more poorly when acting second, only winning 21 out of 56 games for a win rate of 37.5%. Comments by the human players suggested that they thought the expert-rule-based player made good decisions generally, but was a little too cautious in its play so that they were able to win some games they believed they should have lost because the expert-rule-based player did not act as aggressively as he might have done in some situations where he had an advantage.

V. MCTS TREES WITH DETERMINIZATION

MCTS has been applied to a range of games and puzzles and often provides good performance in cases where tree depth/width and difficulty of determining an evaluation function for nonterminal states make depth-limited minimax search ineffective. Modifications are often used to improve basic

MCTS, for example, by ignoring move ordering and using rapid action value estimate (RAVE) values [34] to seed values at previously unexplored nodes which share similarities with already-explored nodes, improved rollout strategies [35], or by using heuristic approaches to limit the number of children for each node [36].

Recent advances in probabilistic planning have presented the idea of *determinization* as a way to solve probabilistic problems [37]. Essentially, each stochastic state transition is determinized (i.e., fixed in advance), and then generates a plan based on the resulting deterministic problem. If the planner arrives at an unexpected state while testing his plan, then he replans using the unexpected state as a starting point and a new set of determinized stochastic state transitions. This approach was extended and generalized by the technique of hindsight optimization [38], which selects among a set of determinized problems by solving determinizations of the future states of a probabilistic problem, resulting after an AI agent's decision state.

MCTS is also making progress in dealing with large partially observable Markov decision problems (POMDPs). Silver and Veness [39] applied MCTS to POMDPs and developed a new algorithm, partially observable Monte Carlo planning (POMCP), which allowed them to deal with problems several orders of magnitude larger than was previously possible. They noted that by using MCTS they had a tool which was better able to deal with two issues that affect classic full width planning algorithms such as value iteration [40]. The *curse of dimensionality* [41] arises because in a problem with n states, value iteration reasons about an n -dimensional belief state. MCTS *samples* the state transitions instead of having to consider them all and so is able to deal with larger state spaces. The *curse of history* [41], in that the number of histories is exponential in the depth, is also dealt with by sampling the histories, and heuristically choosing promising actions using the UCB formula, allowing for a much larger depth to be considered.

Using determinization as a way of dealing with uncertainty is not new. One of the approaches used in the *Bridge* program GIB [21] for playing out the trick taking portion of the game was to select a fixed deal, consistent with bidding, play so far, and find the play resulting in the best expected outcome in the resulting perfect information system. GIB utilized partition search [42] to greatly speed up a minimax/alpha-beta search of each determinized deal, allowing 50 simulations per play on 1990s computer hardware, and ultimately yielding play approaching the standard of human experts. It is interesting to note for GIB that using a relatively small number of determinizations is effective if they are carefully chosen.

Frank and Basin [19] provided a critique of the determinization approach, showing that it is prone to two specific problems that limit the effectiveness of the search. *Strategy fusion* is the problem that different actions are indicated when using determinization from states of the imperfect information game (actually information sets) which are indistinguishable to the player. *Nonlocality* occurs since the values of nodes in an imperfect information tree are affected by decisions higher up the tree, where opponents are able to steer the game toward certain states and away from other (indistinguishable) states; this does not happen for perfect information games or for deter-

minizations. In their work on *Klondike Solitaire*, Bjarnason *et al.* [20] highlighted this issue, providing an example whereby the search would equally favor two moves, where one required foreknowledge of hidden information and another did not. Russell and Norvig called this kind of overoptimism “averaging over clairvoyance” [43], and noted that determinization is incapable of considering issues of information gathering and information hiding. Despite this, perfect information Monte Carlo search (PIMC) has generated strong players in a number of game domains including *Bridge* [21] and *Skat* [18].

A recent study [44] has investigated why PIMC search gives strong results despite its theoretical limitations. By examining the particular qualities of imperfect information games and creating artificial test environments that highlighted these qualities, Long *et al.* [44] were able to show that the potential effectiveness of a PIMC approach was highly dependent on the presence or absence of certain features in the game. They identified three features: *leaf correlation*, *bias*, and *disambiguation*. Leaf correlation measures the probability that all sibling terminal nodes in a tree have the same payoff value; bias measures the probability that the game will favor one of the players; and disambiguation refers to how quickly hidden information is revealed in the course of the game. The study found that PIMC performs poorly in games where the leaf correlation is low, although it is arguable that most sample-based approaches will fail in this case. PIMC also performed poorly when disambiguation was either very high or very low. The effect of bias was small in the examples considered and largely dependent on the leaf correlation value. This correlates well with the observed performance in actual games with PIMC performing well in trick taking games such as *Bridge* [21] and *Skat* [18] where information is revealed progressively as each trick is played so that the disambiguation factor has a moderate value. The low likelihood of the outcome of the game hinging on the last trick also means that leaf correlation is fairly high.

In contrast, *Poker* has a disambiguation factor of 0 as the hidden information (the player's hole cards) is not revealed until the end of the hand. This indicates that PIMC would not perform well at the game. Indeed, recent research in *Poker* has been moving in a different direction, using the technique of counterfactual regret minimization (CFR) [45]. This is a method of computing a strategy profile from the game tree of an extensive form game. It has been shown that for an extensive form game it can be used to determine Nash equilibrium strategies. CFR, and its Monte Carlo sampling-based variant MCCFR [46], is much more efficient than previous methods of solving extensive game trees such as linear programming [47], and has increased the size of game tree that can be analyzed by two orders of magnitude [45], [46]. By collecting *Poker* hands into a manageable number of “buckets” MCCFR can be used to produce strong players for heads up *Texas Hold'Em Poker* [48].

M:TG is a good candidate for investigation by PIMC methods. Leaf correlation in the game is high as it is rare that games are won or lost on the basis of one move at the end of the game; it is more usual for one player to develop the upper hand and apply sufficient continuous pressure on their opponent to win the game. The progressive nature of having an initial hand, unseen by the opponent, drawing cards from an unknown deck

and playing them out into a visible play area also leads to a disambiguation factor that grows slowly throughout the course of the game.

Determinization and MCTS have also been considered for probabilistic planning problems with only one “player.” Bjarnason *et al.* [20] examined the use of UCT in combination with hindsight optimization. They compared using UCT as a method for building determinized problem sets for a hindsight optimization planner and showed that it provided state-of-the-art performance in probabilistic planning domains.

Generating multiple MCTS trees simultaneously in parallel for the same position has also been examined, usually for performance and speed reasons [49], [50]. The idea of searching several independent trees for the same position and combining the results is known as *ensemble UCT* [51], or *root parallelization* in an implementation with concurrency [49], [50], and has been shown in some situations to outperform single-tree UCT given the same total number of simulations [50].

VI. MCTS ALGORITHM DESIGN FOR *M:TG*

In this paper, we combine the methods of ensemble UCT and determinization. We build multiple MCTS trees from the same root node, and for each tree, we determinize chance actions (card draws). Each tree then investigates a possible future from the state space of all possible futures (and the tree of information sets). The determinization of card draws is made as we build the MCTS tree, as late as possible. The first time we reach a state s where we would be required to create chance nodes for a card draw we sample one card draw at random as a specific action a , which takes us to the new state s' ; thereafter, whenever we visit state s in the MCTS tree, we immediately transition to s' without any further sampling; this “lazy determinization” approach is also taken by the HOP-UCT algorithm of Bjarnason *et al.* [20]. As the MCTS tree grows, we effectively fix in place an ordering for the cards in each player’s deck.

If our tree considers all possible outcomes for each chance node in *M:TG*, we may consider this as a single chance node at the top of the tree with an enormous branching factor, or we may branch for each potential card drawn at each chance node. There are 60 cards in a typical *M:TG* deck, and one deck for each player, providing an upper bound of $(60!)^2$ on the number of deals. Since repeat copies of individual cards are allowed (and expected) there will often only be about 15 different cards, and in many games only around 20 cards will be drawn from each deck, but this still yields a combinatorial explosion of possible deals. There are typically 5–6 moves available at a decision node, so this gives a branching factor of approximately 75–90 at 1 ply, around 7000–8000 at 2 ply, and approaching a million at 3 ply. The number of simulations that would be required to generate an MCTS tree capable of collecting meaningful statistics about state values, for all possible states, quickly becomes intractable with increasing depth.

A. Relevance of Individual MCTS Trees

When creating a determinized ordering of cards, as well as being consistent with the game play so far, it seems sensible to try to avoid bias which would make the game an easy win for one of the players. *M:TG* is particularly prone to this, and

indeed this is one of the reasons we believe *M:TG* provides an interesting case study for MCTS research.

We formulate the idea of an “interesting” card ordering: one in which the decisions of the player have an impact on the way the game progresses. We define an ordering as “interesting” if the play “no move” (effectively passing the turn) gives a different result to playing the move suggested by the expert-rule-based player, over a number of game rollouts.

It is not necessarily straightforward to find an “interesting” ordering for a given game state and, indeed, there may not be any ordering of cards that would qualify for our definition of “interesting” if the game state is already heavily biased toward one of the players.

We test whether a card ordering is interesting by generating a random order of cards and carrying out two rollouts from the current game state using that ordering: one with the player making no move and one with the player making the move the expert-rule-based player would have chosen. If the outcome of the game is different between these two rollouts, then the card ordering is classified as “interesting.” We test a small number of random rollouts for each candidate card ordering, and if any one of them yields an “interesting” result, then we accept that card order as interesting. These tests do, of course, computation time and there is a limit to how much time can be sensibly spent searching for an interesting ordering. Ultimately, if we consistently fail to find an interesting ordering, then we must accept that there might not be one to find, at least not within a reasonable time scale. If an interesting ordering is not found, then we use an arbitrarily chosen randomly generated ordering.

An interesting card ordering could be applied to the game at several levels. Preliminary experiments considered using a fraction of the overall simulation budget to: 1) find an interesting ordering for the simulations from each leaf node during MCTS; and 2) find an interesting ordering for the whole deck at the root node only. These were found to give similar, modest improvements in playing strength, but we take option 2 forward since option 1 significantly slows down the search time, by a factor of up to 2, whereas no slowdown is evident for option 2.

Further preliminary experiments were conducted to investigate the budget of overall simulations used to find interesting deck orderings. For the whole tree at the root node, the maximum number of simulations used to find an interesting ordering was varied from 0% to 5%, with good results generally found around the 5% level. This is a maximum and an interesting ordering was usually found in a small fraction of this number of simulations.

A further preliminary investigation looked at whether it was better to use the fixed interesting ordering during simulation rollouts or to revert to the standard random rollouts. These two options were comparable, and random rollouts were chosen in later experiments.

B. Structure of the MCTS Tree

We investigate two families of methods for increasing the effectiveness of search in each determinized MCTS tree.

1) *Dominated Move Pruning*: In building any search tree, limiting the nodes that are added to the tree in order to reduce the scope of the search has often been seen to provide increases




Fig. 3. Potential moves from a position where the player holds cards A, B, and C, with mana costs 4, 3, and 2, respectively, and has five mana available.

in playing strength [52], [36], [35]. In this respect, MCTS is no different from any other tree searching method. How moves are pruned is generally domain dependent. We examined two levels of pruning for our restricted version of *M:TG*, based on incorporating limited heuristic knowledge. The first level of pruning was based around the fact that it is necessary to play land cards before any other cards can be played and that there is little strategic benefit to not playing land when you are able to do so. Nonland pruning prunes any move that does not contain a land card when the player has land in their hand, ensuring that only moves that add more land into the game are considered.

The second, higher, level of pruning makes use of the fact that moves in *M:TG* are frequently composed of multiple cards and that the player chooses a subset of the cards in his hand when he decides on a move. This level of pruning, which we call dominated move pruning, removes any move that is a proper subset of another legal move, so that a maximal set of cards is played.

In summary, the following move pruning strategies were investigated.

- 1) No move pruning. At this level, we consider all possible moves available to each player.
- 2) Nonland pruning. At this level, we prune any move that does not contain a land card if the same move with a land card is available.
- 3) Dominated move pruning. At this level, we prune any move that plays a subset of the cards of another available move.

2) *Binary Decisions*: *M:TG* is unusual among card games in that the moves available on a given turn in the game are a subset of all possible combinations of cards in the player's hand rather than being a single action or a single card. Moreover, the played card group remains active in play rather than being a passive group, as in a melding game such as *Continental Rummy* [53].

Consider that a player has three nonland cards in hand and five land in play. We always suppose here that if land is held, it will be played. Suppose that the cards are A, B, and C, with mana costs of 4, 3, and 2, respectively. The player has five available moves, as shown in Fig. 3.

Here we investigate the case where each node has at most two children, representing the decision to play a card or not. This is illustrated in Fig. 4. With a fixed number of simulations per tree this will substantially increase the depth of the tree, compared to a nonbinary tree which looks the same distance into the future. However, it should allow statistics for good partial decisions (i.e., on whether to play a card) to accumulate independently of




Fig. 4. Binary decision tree of moves corresponding to Fig. 3.

other cards played. Hence, we are able to investigate MCTS decision making on a tree which allows compound moves to be decomposed so that parts of a move can be reinforced separately. This idea, of decomposing a single decision into a sequence of smaller ones as successive levels in the MCTS tree, is similar to the move grouping approach of Childs *et al.* [54].

We imagine this will also be useful in other applications where MCTS is used to choose a subset, for example, we might use this in *M:TG* to select attackers and/or blockers. In this paper, we investigate only the impact upon the decision of which cards to play.

When using this approach, it is desirable that “important” decisions be higher in the binary tree, although it is often difficult to determine *a priori* a sensible importance ordering. Extensive preliminary experiments showed promise for this approach, but did not show any significant difference between using ascending/descending/random orderings based on mana cost. We use descending mana cost in the experiments in Section VII-D, based on the intuition that it will often be stronger to play large creatures first.

C. Simulation Strategies

While MCTS can use approaches to simulation which select randomly among all possible moves, work on MCTS approaches to computer *Go* suggested that using heuristics to guide the simulations provided stronger play [31], but also that a stronger playing strength used for rollouts does not necessarily yield higher playing strength when used in an MCTS framework [34], probably due in large part to the bias that this may introduce.

In our simulation rollouts, we investigate rollouts based on both of our rule-based players. The expert-rule-based player provides a highly structured and completely deterministic rollout and the reduced-rule-based player provides a stochastic approach with some heuristic guidance. We also (briefly) investigated an approach which chose uniformly at random among possible moves during rollouts.

D. Discounted Reward

A player based on MCTS or another forward-looking tree search approach will often make weak moves when in a strong winning (or losing) position. The requirement for the search to

TABLE I
SUMMARY OF EXPERIMENTAL PARAMETERS FOR SECTION VII

Short name	Description	Trees	Simulations per tree	UCT constant	Win/loss reward	Reward discount	Move pruning	Simulation strategy	Tree structure
AP	All Possible Deals / Uniform Random Rollouts	40	250	1.5	1 / 0	1	None	Uniform Random	All Possible Deals
BA	Baseline	40	250	1.5	1 / 0	0.99	Land	Reduced Rules	Unlimited Degree
IL	Interesting Simulations (Leaf. 1% of sim budget)	40	250	1.5	1 / 0	0.99	Land	Reduced Rules / Interesting (Leaf)	Unlimited Degree
IR	Interesting Simulations (Root. 5% of sim budget)	40	250	1.5	1 / 0	0.99	Land	Reduced Rules / Interesting (Root)	Unlimited Degree
NL	Negative Reward for Loss	40	250	1.5	1 / -1	0.99	Land	Reduced Rules	Unlimited Degree
MP	Dominated Move Pruning	40	250	1.5	1 / 0	0.99	Dominated	Reduced Rules	Unlimited Degree
BT	Binary Tree (Descending Mana Cost)	40	250	1.5	1 / 0	0.99	Land	Reduced Rules	Binary

be kept under pressure has been observed repeatedly [55]. In order to create a sense of urgency within the player we use an idea from many game tree search implementations (e.g., [30]) and discount the reward value that is propagated back up the tree from the terminal state of a simulation. If the base reward is γ and it takes t turns (counting turns for both players) to reach a terminal state from the current root state, then the actual reward propagated back through the tree is $\gamma\lambda^t$ for some discount parameter λ with $0 < \lambda \leq 1$. Here we choose $\lambda = 0.99$ which yields discount factors between 0.7 and 0.5 for a typical *M:TG* game of 40 to 60 turns.

We also compare the effects of assigning a loss reward of 0 or -1 (a win having a reward of $+1$ in both cases). The value of -1 , in combination with discounted rewards, aims to incentivize the player to put off losses for as long as possible. This can be beneficial, as extending the length of the game increases the chance of obtaining a lucky card draw.

VII. EXPERIMENTS

Our empirical investigation compares MCTS players for *M:TG* using the approaches explained in Section VI (using parameters from Table I). In Section VII-A, we present a simple experiment to show that a naïve implementation of UCT does not yield strong play. In Section VII-B, we explore the effect of varying the number of determinizations for a fixed simulation budget, and show that with 10 000 simulations, around 40 determinizations, each with 250 simulations, provides good play (a similar result was found for the card game *Dou Di Zhu* in [56]). In Section VII-C, we compare the relative performance of the approaches in Table I. In Section VII-D, we evaluate the effectiveness of combinations of approaches. The baseline conditions reported in Table I are a result of extensive preliminary experiments (some of which are reported in Section VII-B).

The cards that comprise the deck used by the players are fixed in advance and both players utilize the same deck composition. We created a selection of *M:TG* style creature and land cards for the decks. The decks contain 40 cards with 17 land cards and 23 creature cards. These proportions are the same as ones generally

used by competitive *M:TG* players in tournaments, as they represent the perceived wisdom of providing the best probability to draw a useful mix of land and spells throughout the game. The 23 creatures in the deck were spread among a range of combinations of power, toughness, and cost from $1/1^{(1)}$ to $6/6^{(7)}$.

To provide consistency between experiments, and reduce the variance of our results, in the experiments in Sections VII-A and VII-B, we randomly generated and tested fixed deck orderings until we had 50 orderings that were not particularly biased toward either of the players. In Section VII-C, we use 100 unbiased fixed orderings for each pair of players. This type of approach is used in a variety of games to reduce the variance between trials, and notably used in *Bridge* and *Whist* tournaments [57] between high-level human players. The experiments were carried out twice with the players alternating between player 1 and player 2 positions, to further reduce bias due to any advantage in going first/second.

Our experiments were carried out on a range of server machines. Broadly speaking we wanted to maintain decision times of around 1 computation time second or less, since that would be acceptable in play versus a human player. We use number of simulations as the stopping criterion in all cases. Computation times are reported for a server with an Intel Xeon X5460 processor, and 4-GB RAM, running Windows Server 2003. Code was written in C# for the Microsoft .NET framework.

A. MCTS for All Possible Deals

As remarked earlier, the branching factor at a chance node involving a single card draw may be 15 or higher, and since there is a chance node for each decision node in *M:TG*, this doubles the depth of the tree compared to determinization approaches, which fix these chance nodes in advance. While we would not expect MCTS to perform well for a tree which grows so rapidly with depth, it provides an important baseline for our experiments. The approach is illustrated in Fig. 5. Note that in this case as well as other experiments (unless stated otherwise), card draws were only specified at the last possible moment (i.e., at the point of drawing a card).




Fig. 5. An MCTS tree with chance nodes.

There are multiple methods that can be used in order to select a chance node when descending the tree. Here we select chance outcomes (card draws) uniformly at random. However, since in practice there are repeated cards in the deck, actually we only have one chance outcome per card type, and weight this according to the number of cards of that type in the deck. Another possibility, not considered here, is that one of the players in the game chooses the card to be drawn, with the active player selecting the “best” chance action and the nonactive player selecting the “worst” chance action. In all of these cases, the number of nodes in the tree increases very rapidly with each chance node, which is likely to lead to poor playing strength for MCTS.

The all possible deals player was played against the expert-rule-based player and the reduced-rule-based player, using simulation rollouts that select uniformly at random from among the available legal moves. Over ten replications of 100 games we see that the all possible deals player is significantly weaker than the expert-rule-based or reduced-rule-based players, winning only 23% of games against the expert-rule-based player and 38% of games against the reduced-rule-based player. This result provides a baseline to which we can compare our other experimental results in order to determine if our adjustments to the MCTS algorithm are having a beneficial effect.

B. Varying the Number of Determinizations

When using determinization, for a fixed budget on the total number of simulations, we trade off the number of determinization trees versus the number of simulations per tree. If the number of determinizations is too low, we may get a poor result since the small sample of determinizations is not representative of the combinatorially large set of deck orderings. If the number of simulations per tree is too small, then MCTS has not had enough time to exploit promising play lines in the tree for each determinization. We run the tests for a fixed number of total simulations on each tree and then simply add the results from all the trees together and select the move that has the most number of visits over all trees.

In Table II and Fig. 6, we vary the number n of determinizations, with each determinization tree having around $10\,000/n$ simulations. Other experimental conditions are as for the baseline player in Table I.

The first thing we note in Table II is that using an ensemble of determinizations yields much stronger play than the naïve




Fig. 6. Comparison of the effect of using multiple determinized trees.

TABLE II
WIN RATE (IN PERCENT) OF MCTS PLAYER WITH MULTIPLE DETERMINIZED TREES AND 10 000 SIMULATIONS IN TOTAL

No. of trees	Expert Rules Simulations		Reduced Rules Simulations	
	vs Reduced Rules	vs Expert Rules	vs Reduced Rules	vs Expert Rules
1	27	26	31	32
2	31	29	44	38
4	45	34	48	36
5	46	35	38	40
10	45	33	44	40
20	41	36	46	44
40	48	31	56	47
50	42	34	52	46
100	55	33	56	43
250	50	29	45	47
500	38	21	48	46
1000	15	12	41	26

MCTS implementation in Section VII-A. We see also that using reduced rules simulations gives better results than using expert rules simulations, even though the reduced-rule-based player is much weaker than the expert-rule-based player. It seems the reduced-rule-based player provides enough focus to make simulation results meaningful for trees of this size (compared with the results of Section VII-A) while not rigidly defining game outcomes (as for the expert-rule-based player). Similar results are reported for *Go* in [34]. In Sections VII-C and VII-D, we will consider only these more effective reduced rules simulations.

In each case, we see that the best number of determinizations occurs between 20 and 100, and the best number of simulations per determinization between 500 and 100, with a total budget of 10 000 simulations. This and results from [56] motivate us to choose 40 determinizations with 250 simulations per determinization tree in Sections VII-C and VII-D.

The computation time used for a single move decision increases slightly as the number of trees increases, from 0.62 s for a single tree to 1.12 s for 50 trees. Inefficiencies in our code (and particularly the way in which trees are combined) increase the computation time per move up to 14.01 s per move for 1000 trees, although this could be significantly reduced below 1 s per move through more careful design.

Similar experiments were conducted with a budget of 100 000 simulations and the number of determinizations n taking values from the set $\{1, 2, 4, 5, 10, 20, 50, 100, 500, 1000, 2000, 5000, 10\,000\}$, with about $100\,000/n$ simulations per determinization.




Fig. 7. Average win rate (in percent) of players in round robin tournament. Error bars show 95% confidence intervals.

The best number of simulations per determinization again lay in the range from 100 to 1000, suggesting that an increased simulation budget is best used in running additional determinizations rather than searching each determinization more deeply. The effects on playing strength of more simulations are analyzed in Table VI.

C. Comparison of MCTS Enhancements

We have outlined a number of different enhancements to MCTS, all with the potential for improving the performance of the search when utilized in a game such as *M:TG*. A round robin tournament was conducted with representative players from each approach (as shown in Table I), to provide a measure of comparative strength of the various enhancements.

In the tournament, each player played each other player over 100 games with 50 games being played as each of player 1 and player 2. The same, fixed 50 deck orderings are used for each match, to minimize variance and provide a fair comparison. The results are shown in Table III, with average win rates for each player in Fig. 7. Each player used a fixed budget of 10 000 simulations; Table IV shows average computation times per decision, from which we can generally see that BA, IR, NL, and MP approaches take approximately the same amount of computation time. The AP approach is slower, due to the overhead of generating a much wider tree than other approaches. The IL approach, which consumes extra time at every leaf node in the tree searching for an “interesting” determinization, is understandably by far the slowest method. The low branching factor of the BT approach leads to a much lower average time per move.

The benefits of ensemble determinization are clear, with all other players greatly outperforming the all possible deals (AP) player which attempts to construct the whole tree without the focusing effect of determinization. All of our enhancements to the basic ensemble determinization approach (IL, IR, NL, MP and BT) improve on the baseline (BA) approach, with the difference significant at the 95% level for all except for negative reward for loss (NL). Methods which maintain “pressure” on the MCTS, either by finding “interesting” determinizations (IL, IR) or by rewarding delaying tactics when behind (NL) are seen to enhance performance over the baseline player. The use of domain knowledge to prune the tree (MP) is also seen to be effective when compared to the baseline.




Fig. 8. A binary MCTS tree.

TABLE III
WIN RATE (IN PERCENT) OF MCTS PLAYERS IN ROUND ROBIN
TOURNAMENT WIN RATE FOR ROW PLAYER

Player	AP	BA	IL	IR	NL	MP	BT
AP		23	26	21	26	25	18
BA	67		42	40	44	41	44
IL	74	58		46	48	53	48
IR	79	60	54		58	53	43
NL	74	56	52	42		45	44
MP	75	59	47	47	55		52
BT	82	56	52	57	56	48	

TABLE IV
AVERAGE COMPUTATION TIME PER MOVE FOR THE MCTS PLAYERS IN
SECTION VII-C

Player	Average time per move (seconds)
AP	5.65
BA	0.75
IL	9.81
IR	1.00
NL	1.07
MP	1.00
BT	0.23

The IL, IR, MP, and BT approaches have similar playing strength, with BT and IR slightly in front, although not significantly so. These four approaches are quite different in the way that they enhance the baseline algorithm. The fact that they enhance different aspects of the ensemble determinization approach is further evidenced by their nontransitive performance against each other. For example, the BT approach beats the otherwise unbeaten IR approach, and IR beats MP, but MP is stronger than BT.

The interesting simulations (root) (IR) result is slightly better than the interesting simulations (leaf) (IL) result, although IR consumes significantly less computation time than IL for a given number of simulations (1.00 s per decision for IR versus 9.81 s for IL; see Table IV). Hence, we have evidence in support of the




Fig. 9. A nonbinary MCTS tree.

effectiveness of finding interesting determinizations, but it does not appear that we need the detail or computational expense of attempting to find an interesting simulation at every leaf of the tree. This observation leads us to use the IR variant in the combination experiments in Section VII-D.

The use of binary trees (BT) is consistently strong against all players, losing only to the dominated move pruning (MP) player. This is particularly notable since the approach is more than three times as fast as any other approach. Figs. 8 and 9 illustrate the difference in tree structure for the binary tree enhancement. We believe that the idea of using binary trees, in combination with domain knowledge, will likely lead to further enhancements, and begin the exploration of this in Section VII-D. However, due to the difficulty in finding appropriate domain knowledge, this is a large piece of work in itself, and we anticipate future work in this area.

D. Combinations of MCTS Enhancements

We have shown in Section VII-C that our enhancements to the basic MCTS algorithm individually produce a stronger player than the baseline MCTS approach to using ensemble determinization. In this section, we investigate the effectiveness of combinations of some of the best performing enhancements. We took four enhancements that had performed strongly in individual experiments and tested all possible combinations of them. The enhancements tested were as follows (in each case the expected stronger level is listed first).

- Binary trees (BT/–): used (ordered by descending mana cost)/not used.
- Move pruning (MP/–): dominated move pruning/nonland pruning.
- Negative reward for loss (NL/–): –1 reward for loss/0 reward for loss.
- Interesting simulations (root) (R/–): used (at most 5% of simulation budget used to find an interesting ordering)/not used.

In the results of this section, we denote each player as a four-tuple, to denote the level of each enhancement. For example, the player (BT, MP, –, –) utilizes binary trees and dominated move pruning, but not the negative reward for loss or interesting simulations (root). These experiments are very computation time intensive, and 100 replications were conducted using a large computation time cluster.

We present average performance versus the expert-rule-based player in Table V. In addition to the results given in the table,

TABLE V
COMBINATION EXPERIMENTS—AVERAGE WIN RATE
(IN PERCENT) OVER 100 TRIALS

Player	Win % vs Expert Rules Player	Average time per move (seconds)
(BT, MP, NL, IR)	49.5	0.21
(BT, MP, NL, –)	50.5	0.17
(BT, MP, –, IR)	50.4	0.27
(BT, MP, –, –)	50.5	0.20
(BT, –, NL, IR)	47.0	0.31
(BT, –, NL, –)	47.7	0.23
(BT, –, –, IR)	47.6	0.28
(BT, –, –, –)	47.6	0.19
(–, MP, NL, IR)	47.4	1.01
(–, MP, NL, –)	44.5	1.05
(–, MP, –, IR)	47.7	1.00
(–, MP, –, –)	46.1	1.05
(–, –, NL, IR)	48.3	1.00
(–, –, NL, –)	43.6	0.92
(–, –, –, IR)	47.7	1.01
(–, –, –, –)	43.9	0.79

we observed that reduced rules rollouts significantly outperform expert rules rollouts (by around 10% in most cases), and that all the players which use at least one enhancement significantly outperform the reduced-rule-based player.

The results were analyzed using multiway analysis of variance (ANOVA) [58], using the *R* statistical package [59]. Multiway ANOVA showed that enhancements (BT, MP, and IR) yielded performance improvements which were significant at the 99% level [i.e., that (BT, *, *, *) significantly outperforms (–, *, *, *), etc.]. NL represented a significant improvement only at the 90% level. The following pairs of enhancements were also significant at the 99% level: BT:MP, BT:IR, and MP:IR. Only one triple of enhancements yielded significantly better performance at the 99% level: BT:MP:IR.

ANOVA analysis and the results in Table V show that our proposed enhancements do indeed improve performance of ensemble determined MCTS, in combination as well as individually. The (BT, MP, *, *) players provide the strongest performance, yielding playing strength slightly better than the expert-rule-based player. Achievement of a higher than 50% win rate is a substantive achievement when we consider the strength of the expert-rule-based player against expert human opponents, and the fact that the (BT, MP, *, *) players achieve this performance without using the knowledge encoded in these expert rules.

The BT enhancement significantly decreases the computation time per decision, probably as a result of the MCTS selection phase having far fewer branches to choose between at

TABLE VI
A HUNDRED THOUSAND ROLLOUTS COMBINATION EXPERIMENTS
—AVERAGE WIN RATE (IN PERCENT) OVER 40 TRIALS

Player	Win % vs Expert Rules Player	Average time per move (seconds)
(BT, MP, IR)	51.2	3.21
(BT, MP, —)	51.5	3.21
(BT, —, IR)	51.6	3.63
(BT, —, —)	52.0	3.36
(—, MP, IR)	44.3	8.38
(—, MP, —)	44.7	7.91
(—, —, IR)	47.2	6.83
(—, —, —)	37.3	6.45

each level in the tree. MP yields a slight improvement in computation time when coupled with BT. The other enhancements slightly increase the computation time per decision, but not significantly so.

The results of our analysis underline the utility of all the proposed methods, the dominance of the BT:MP combination, and the complexity of the interaction between methods in yielding increased playing strength.

We carried out additional experiments in order to investigate whether increasing the number of rollouts to 100 000 would provide any significant increase in the performance of the most promising combinations. In this case, we did not consider the negative reward for loss (NL) enhancement (using a reward for loss of zero) due to the computation-time-intensive nature of these experiments and the fact that the previous results suggest that it was the least effective of the four enhancements. The results of this are shown in Table VI. Note that these experiments are very time consuming, requiring roughly five to ten times as much computation time per trial as those in Table V.

We see here modest improvements in overall performance, when using the BT enhancement with or without other enhancements. Counterintuitively, without this enhancement, performance is no better and indeed slightly worse than when using a smaller simulation budget. We have observed this phenomenon for other games of partial information [56], [60] which probably arises due to the large branching factor as we descend the tree even when determinization is used, so that the additional simulation budget is used in chasing somewhat arbitrary decision possibilities. That BT mitigates this problem suggests that this is a particularly interesting area for further study, capable of focussing search into interesting areas of the tree. BT likely improves matters here since the reduction of the degree of the tree results in a more focussed search in each determinization.

VIII. CONCLUSION

In this paper, we have introduced the popular card game *M:TG*. We believe *M:TG* is an interesting domain for computational intelligence and AI, and particularly MCTS, for a variety of reasons. The game is highly popular and commercially successful, and has (human) players at professional levels. It is an imperfect information game, with unique cards that provide a rich level of tactical play and provide a very high branching factor for any search-based approach. Expert heuristics are difficult to formulate because of the variety and complexity of

the game situations that arise and the fact that the effectiveness of many actions is highly dependent on the current game state. All of these factors suggest that *M:TG* would be an extremely difficult challenge for conventional evaluation-based search methods.

We also feel that the structure of the game is suited to analysis by MCTS. The progressive revealing of information as players draw new cards from their decks and play them out, combined with the relative unlikelihood of similar game states leading to radically different game outcomes, are both features that suggest that MCTS should be able to generate strong play.

The central theme of this paper is the use of multiple determinized trees as a means of dealing with imperfect information in an MCTS search. We have shown that this approach provides significant benefits in playing strength, becoming competitive with a sophisticated expert-rule-based player with a simulation budget of less than one computation time second on standard hardware, despite having no access to expert knowledge. In addition to that, we have presented a wide variety of enhancements to the determinized trees and analyzed the effect on playing strength that each enhancement offers. All of these enhancements show further improvement. We investigated a modification of the structure of the decision tree to a binary tree, well suited to *M:TG* where decisions amount to the choice of a subset of cards from a small set, rather than an individual card. As well as providing significant improvements in playing strength, the binary tree representation substantially reduced computation time per move. Dominated move pruning used limited domain knowledge, of a type applicable to a wide variety of games involving subset choice, to significantly reduce the branching factor within the tree. Another promising approach maintained pressure on the MCTS algorithm by choosing “interesting” determinizations which were balanced between the two players. An enhancement which used decaying reward to encourage delaying moves when behind had some positive effect, but was not as effective as the preceding three enhancements.

The rollout strategy had a profound effect in our experiments. Applying a fully deterministic rollout strategy, as we did when using our expert-rule-based player to handle the rollouts, provided a clearly inferior performance to utilizing the reduced-rule-based player which uses very limited domain knowledge, but incorporates some randomness within its decisions. This was true in all of our experiments and despite the fact that the expert-rule-based player is an intrinsically stronger player than the reduced-rule-based player. However, using a naïve rollout strategy which chose uniformly at random from all possible moves proved to be very weak.

MCTS, suitably enhanced by the range of approaches we have suggested in this paper, was able to compete with, and outperform, a strong expert-rule-based player (which is in turn competitive with strong human players). Hence, the paper adds to the volume of work which suggests MCTS as a powerful algorithm for game AI, for a game of a somewhat different nature to those previously studied.

In future work, we will look at increasing the complexity of the game environment by including a wider variety of *M:TG* cards and card types. This will increase the scope of the tactical decisions available to the player and will make it significantly

harder to encode strong knowledge-based players. We also intend to look more closely at binary trees in conjunction with domain knowledge, which we believe may yield significant further improvements in playing strength.

Card and board games such as *MTG* provide excellent test beds for new artificial intelligence and computational intelligence techniques, having intermediate complexity between perfect information games such as *Chess* and *Go*, and video games. As such we believe they represent an important stepping stone toward better AI in commercial video games.

ACKNOWLEDGMENT

The authors would like to thank the anonymous reviewers for their helpful comments.

REFERENCES

- [1] H. J. van den Herik, "The drosophila revisited," *Int. Comput. Games Assoc. J.*, vol. 33, no. 2, pp. 65–66, 2010.
- [2] C.-S. Lee, M.-H. Wang, G. M. J.-B. Chaslot, J.-B. Hoock, A. Rimmel, O. Teytaud, S.-R. Tsai, S.-C. Hsu, and T.-P. Hong, "The computational intelligence of MoGo revealed in Taiwan's computer Go tournaments," *IEEE Trans. Comput. Intell. AI Games*, vol. 1, no. 1, pp. 73–89, Mar. 2009.
- [3] A. Rimmel, O. Teytaud, C.-S. Lee, S.-J. Yen, M.-H. Wang, and S.-R. Tsai, "Current frontiers in computer Go," *IEEE Trans. Comput. Intell. AI Games*, vol. 2, no. 4, pp. 229–238, Dec. 2010.
- [4] C.-S. Lee, M. Müller, and O. Teytaud, "Guest editorial: Special issue on Monte Carlo techniques and computer Go," *IEEE Trans. Comput. Intell. AI Games*, vol. 2, no. 4, pp. 225–228, Dec. 2010.
- [5] C. Browne, E. J. Powley, D. Whitehouse, S. M. Lucas, P. I. Cowling, P. Rohlfshagen, S. Tavener, D. Perez, S. Samothrakis, and S. Colton, "A survey of Monte Carlo tree search methods," *IEEE Trans. Comput. Intell. AI Games*, vol. 4, no. 1, pp. 1–43, Mar. 2012.
- [6] B. Arneson, R. B. Hayward, and P. Henderson, "Monte Carlo tree search in Hex," *IEEE Trans. Comput. Intell. AI Games*, vol. 2, no. 4, pp. 251–258, Dec. 2010.
- [7] F. Teytaud and O. Teytaud, "Creating an upper-confidence-tree program for Havannah," in *Advances in Computer Games*, ser. Lecture Notes in Computer Science. Berlin, Germany: Springer-Verlag, 2010, vol. 6048, pp. 65–74.
- [8] J. Méhat and T. Cazenave, "Combining UCT and nested Monte Carlo search for single-player general game playing," *IEEE Trans. Comput. Intell. AI Games*, vol. 2, no. 4, pp. 271–277, Dec. 2010.
- [9] Y. Björnsson and H. Finnsson, "Cadiaplayer: A simulation-based general game player," *IEEE Trans. Comput. Intell. AI Games*, vol. 1, no. 1, pp. 4–15, Mar. 2009.
- [10] M. Enzenberger, M. Müller, B. Arneson, and R. B. Segal, "Fuego—An open-source framework for board games and Go engine based on Monte Carlo tree search," *IEEE Trans. Comput. Intell. AI Games*, vol. 2, no. 4, pp. 259–270, Dec. 2010.
- [11] J. Schaeffer, "The games computers (and people) play tree search," *Adv. Comput.*, vol. 52, pp. 189–266, 2000.
- [12] S. E. Siwek, "Video Games in the 21st Century: The 2010 Report," 2010 [Online]. Available: http://www.theesa.com/facts/pdfs/VideoGames21stCentury_2010.pdf
- [13] N. Ersotelos and F. Dong, "Building highly realistic facial modeling and animation: A survey," *Vis. Comput.*, vol. 24, no. 1, pp. 13–30, 2008.
- [14] P. Tozour, "The perils of AI scripting," in *AI Game Programming Wisdom*, S. Rabin, Ed. Newton Center, MA: Charles River Media, 2002, pp. 541–547.
- [15] I. Szita, G. M. J.-B. Chaslot, and P. Spronck, "Monte-Carlo tree search in Settlers of Catan," in *Advances in Computer Games*, ser. Lecture Notes in Computer Science. Berlin, Germany: Springer-Verlag, 2010, vol. 6048, pp. 21–32.
- [16] N. R. Sturtevant, "An analysis of UCT in multi-player games," in *Proceedings of the 6th International Conference on Computers and Games*, ser. Lecture Notes in Computer Science. Berlin, Germany: Springer-Verlag, 2008, vol. 5131, pp. 37–49.
- [17] G. van den Broeck, K. Driessens, and J. Ramon, "Monte-Carlo tree search in poker using expected reward distributions," in *Advances in Machine Learning*, ser. Lecture Notes in Computer Science. Berlin, Germany: Springer-Verlag, 2009, vol. 5828, pp. 367–381.
- [18] M. Buro, J. R. Long, T. Furtak, and N. R. Sturtevant, "Improving state evaluation, inference, and search in trick-based card games," in *Proc. 21st Int. Joint Conf. Artif. Intell.*, Pasadena, CA, 2009, pp. 1407–1413.
- [19] I. Frank and D. Basin, "Search in games with incomplete information: A case study using bridge card play," *Artif. Intell.*, vol. 100, no. 1–2, pp. 87–123, 1998.
- [20] R. Bjarnason, A. Fern, and P. Tadepalli, "Lower bounding Klondike solitaire with Monte-Carlo planning," in *Proc. 19th Int. Conf. Autom. Planning Scheduling*, Thessaloniki, Greece, 2009, pp. 26–33.
- [21] M. L. Ginsberg, "Gib: Imperfect information in a computationally challenging game," *J. Artif. Intell. Res.*, vol. 14, pp. 303–358, 2001.
- [22] Magic: The Gathering, Wizards of the Coast, [Online]. Available: <http://www.magicthegathering.com>
- [23] H. Rifkind, "Magic: Game that made Monopoly disappear," *The Times*, Jul. 2005 [Online]. Available: http://www.timesonline.co.uk/tol/life_and_style/article545389.ece?token=null&offset=0&page=1
- [24] G. Giles, "House of cards," 1995 [Online]. Available: <http://www.metroactive.com/papers/sonoma/11.09.95/magic.html>
- [25] P. Buckland, "Duels of the Planeswalkers: All about AI," 2009 [Online]. Available: <http://www.wizards.com/Magic/Magazine/Article.aspx?x=mtg/daily/feature/44>
- [26] Z. Mowshowitz, "Review and analysis: Duels of the Planeswalkers," 2009 [Online]. Available: <http://www.top8magic.com/2009/06/review-and-analysis-duels-of-the-planeswalkers-by-zvi-mowshowitz/>
- [27] P. Auer, N. Cesa-Bianchi, and P. Fischer, "Finite-time analysis of the multiarmed bandit problem," *Mach. Learn.*, vol. 47, no. 2, pp. 235–256, 2002.
- [28] G. M. J.-B. Chaslot, J.-T. Saito, B. Bouzy, J. W. H. M. Uiterwijk, and H. J. van den Herik, "Monte-Carlo strategies for computer Go," in *Proc. BeNeLux Conf. Artif. Intell.*, Namur, Belgium, 2006, pp. 83–91.
- [29] R. Coulom, "Efficient selectivity and backup operators in Monte-Carlo tree search," in *Proc. 5th Int. Conf. Comput. Games*, Turin, Italy, 2006, pp. 72–83.
- [30] L. Kocsis and C. Szepesvári, "Bandit based Monte-Carlo planning," in *Proc. Euro. Conf. Mach. Learn.*, J. Fürnkranz, T. Scheffer, and M. Spiliopoulou, Eds., Berlin, Germany, 2006, pp. 282–293.
- [31] Y. Wang and S. Gelly, "Modifications of UCT and sequence-like simulations for Monte-Carlo Go," in *Proc. IEEE Symp. Comput. Intell. Games*, Honolulu, HI, 2007, pp. 175–182.
- [32] R. E. Korf, "Depth-first iterative-deepening: An optimal admissible tree search," *Artif. Intell.*, vol. 27, no. 1, pp. 97–109, 1985.
- [33] J. Ferraiolo, "The MODO Fiasco: Corporate Hubris and Magic Online," 2004 [Online]. Available: <http://www.starcitygames.com/php/news/article/6985.html>
- [34] S. Gelly and D. Silver, "Combining online and offline knowledge in UCT," in *Proc. 24th Annu. Int. Conf. Mach. Learn.*, Corvallis, OR, 2007, pp. 273–280.
- [35] G. M. J.-B. Chaslot, C. Fiter, J.-B. Hoock, A. Rimmel, and O. Teytaud, "Adding expert knowledge and exploration in Monte-Carlo tree search," in *Advances in Computer Games*, ser. Lecture Notes in Computer Science. Berlin, Germany: Springer-Verlag, 2010, vol. 6048, pp. 1–13.
- [36] G. M. J.-B. Chaslot, M. H. M. Winands, H. J. van den Herik, J. W. H. M. Uiterwijk, and B. Bouzy, "Progressive strategies for Monte-Carlo tree search," *New Math. Nat. Comput.*, vol. 4, no. 3, pp. 343–357, 2008.
- [37] S. Yoon, A. Fern, and R. L. Givan, "FF-replan: A baseline for probabilistic planning," in *Proc. 17th Int. Conf. Autom. Planning Scheduling*, 2007, pp. 352–359.
- [38] S. Yoon, A. Fern, R. L. Givan, and S. Kambhampati, "Probabilistic planning via determinization in hindsight," in *Proc. Assoc. Adv. Artif. Intell.*, Chicago, IL, 2008, pp. 1010–1017.
- [39] D. Silver and J. Veness, "Monte-Carlo planning in large POMDPs," in *Proc. Neural Inf. Process. Syst.*, Vancouver, BC, Canada, 2010, pp. 1–9.
- [40] L. P. Kaelbling, M. L. Littman, and A. R. Cassandra, "Planning and acting in partially observable stochastic domains," *Artif. Intell.*, 101, no. 1–2, pp. 99–134, 1998.
- [41] J. Pineau, G. Gordon, and S. Thrun, "Anytime point-based approximations for large POMDPs," *J. Artif. Intell. Res.*, vol. 27, pp. 335–380, 2006.

- [42] M. L. Ginsberg, "Partition search," in *Proc. 13th Nat. Conf. Artif. Intell./8th Innovative Appl. Artif. Intell. Conf.*, 1996, pp. 228–233.
- [43] S. J. Russell and P. Norvig, *Artificial Intelligence: A Modern Approach*, 3rd ed. Upper Saddle River, NJ: Prentice-Hall, 2009.
- [44] J. R. Long, N. R. Sturtevant, M. Buro, and T. Furtak, "Understanding the success of perfect information Monte Carlo sampling in game tree search," in *Proc. Assoc. Adv. Artif. Intell.*, Atlanta, GA, 2010, pp. 134–140.
- [45] M. Zinkevich, M. Johanson, M. Bowling, and C. Piccione, "Regret minimization in games with incomplete information," in *Proc. Adv. Neural Inf. Process. Syst.*, Vancouver, BC, Canada, 2008, pp. 1729–1736.
- [46] M. Lanctot, K. Waugh, M. Zinkevich, and M. Bowling, "Monte Carlo sampling for regret minimization in extensive games," in *Proc. Adv. Neural Inf. Process. Syst.*, Vancouver, BC, Canada, 2009, pp. 1078–1086.
- [47] D. Koller and N. Megiddo, "The complexity of two-person zero-sum games in extensive form," *Games Econ. Behav.*, vol. 4, pp. 528–552, 1992.
- [48] M. Bowling, N. A. Risk, N. Bard, D. Billings, N. Burch, J. Davidson, J. Hawkin, R. Holte, M. Johanson, M. Kan, B. Paradis, J. Schaeffer, D. Schnizlein, D. Szafron, K. Waugh, and M. Zinkevich, "Monte Carlo sampling for regret minimization in extensive games," in *Proc. Adv. Neural Inf. Process. Syst.*, 2009, pp. 1391–1392.
- [49] T. Cazenave and N. Jouandeau, "On the parallelization of UCT," in *Proc. Comput. Games Workshop*, Amsterdam, The Netherlands, 2007, pp. 93–101.
- [50] G. M. J.-B. Chaslot, M. H. M. Winands, and H. J. van den Herik, "Parallel Monte-Carlo tree search," in *Proceedings of the 6th International Conference on Computers and Games*, ser. Lecture Notes in Computer Science. Berlin, Germany: Springer-Verlag, 2008, vol. 5131, pp. 60–71.
- [51] A. Fern and P. Lewis, "Ensemble Monte-Carlo planning: An empirical study," in *Proc. 21st Int. Conf. Autom. Planning Scheduling*, Freiburg, Germany, 2011, pp. 58–65.
- [52] D. E. Knuth and R. W. Moore, "An analysis of alpha-beta pruning," *Artif. Intell.*, vol. 6, no. 4, pp. 293–326, 1975.
- [53] Pagat, "Rummy," [Online]. Available: <http://www.pagat.com/rummy/rummy.html>
- [54] B. E. Childs, J. H. Brodeur, and L. Kocsis, "Transpositions and move groups in Monte Carlo tree search," in *Proc. IEEE Symp. Comput. Intell. Games*, Perth, Australia, 2008, pp. 389–395.
- [55] I. Althöfer, "On the laziness of Monte-Carlo game tree search in non-tight situations," Friedrich-Schiller Univ., Jena, Jena, Germany, 2008.
- [56] E. J. Powley, D. Whitehouse, and P. I. Cowling, "Determinization in Monte-Carlo tree search for the card game Dou Di Zhu," in *Proc. Artif. Intell. Simul. Behav.*, York, U.K., 2011, pp. 17–24.
- [57] World Bridge Federation, "General conditions of contest," 2011 [Online]. Available: <http://www.worldbridge.org/departments/rules/GeneralConditionsOfContest2011.pdf>
- [58] M. H. Kutner, J. Neter, C. J. Nachtsheim, and W. Wasserman, *Applied Linear Statistical Models*, 5th ed. New York: McGraw-Hill, 2004.
- [59] Foundation for Statistical Computing, Development Core Team, "A language and environment for statistical computing," 2008.
- [60] P. I. Cowling, E. J. Powley, and D. Whitehouse, "Information set Monte Carlo tree search," *IEEE Trans. Comput. Intell. AI Games*, vol. 4, no. 2, pp. 120–143, 2012.



Peter I. Cowling (M'05) received the M.A. and D.Phil. degrees from Corpus Christi College, University of Oxford, Oxford, U.K., in 1989 and 1997, respectively.


He is a Professor of Computer Science and Associate Dean (Research and Knowledge Transfer) at the University of Bradford, Bradford, U.K., where he leads the Artificial Intelligence Research Centre. In September 2012, he will take up an Anniversary Chair at the University of York, York, U.K., joined between the Department of Computer Science and the York Management School. His work centers on computerized decision making in games, scheduling, and resource-constrained optimization, where real-world situations can be modeled as constrained search problems in large directed graphs. He has a particular interest in general-purpose approaches such as hyperheuristics (where he is a pioneer) and Monte Carlo tree search (especially the application to games with stochastic outcomes and complete information). He has worked with a wide range of industrial partners, developing commercially successful systems for steel scheduling, mobile workforce planning, and staff timetabling. He is a director of two research spinout companies. He has published over 80 scientific papers in high-quality journals and conferences.

Prof. Cowling is a founding Associate Editor of the IEEE TRANSACTIONS ON COMPUTATIONAL INTELLIGENCE AND AI FOR GAMES. He has won a range of academic prizes and "best paper" awards, and given invited talks at a wide range of universities and conference meetings.



Colin D. Ward (S'09) received the B.Sc. degree in computer and information systems from the University of Bradford, Bradford, U.K., in 2008, where he is currently working toward the Ph.D. degree in artificial intelligence. His dissertation examines game domains with incomplete information and search methods that can be applied to those domains, particularly Monte Carlo and Monte Carlo tree search methods.

With his background in computer science and as a competitive game player (he has been ranked among the top 150 *Magic: The Gathering* players in the United Kingdom), his research interests are focused on artificial intelligence and machine learning approaches to decision making in games.



Edward J. Powley (M'10) received the M.Math. degree in mathematics and computer science and the Ph.D. degree in computer science from the University of York, York, U.K., in 2006 and 2010, respectively.

He is currently a Research Fellow at the University of Bradford, Bradford, U.K., where he is a member of the Artificial Intelligence Research Centre in the School of Computing, Informatics and Media. His current work involves investigating Monte Carlo tree search (MCTS) for games with hidden information and stochastic outcomes. His other research interests include cellular automata, and game theory for security.

Dr. Powley was awarded the P B Kennedy Prize and the BAE Systems ATC Prize.